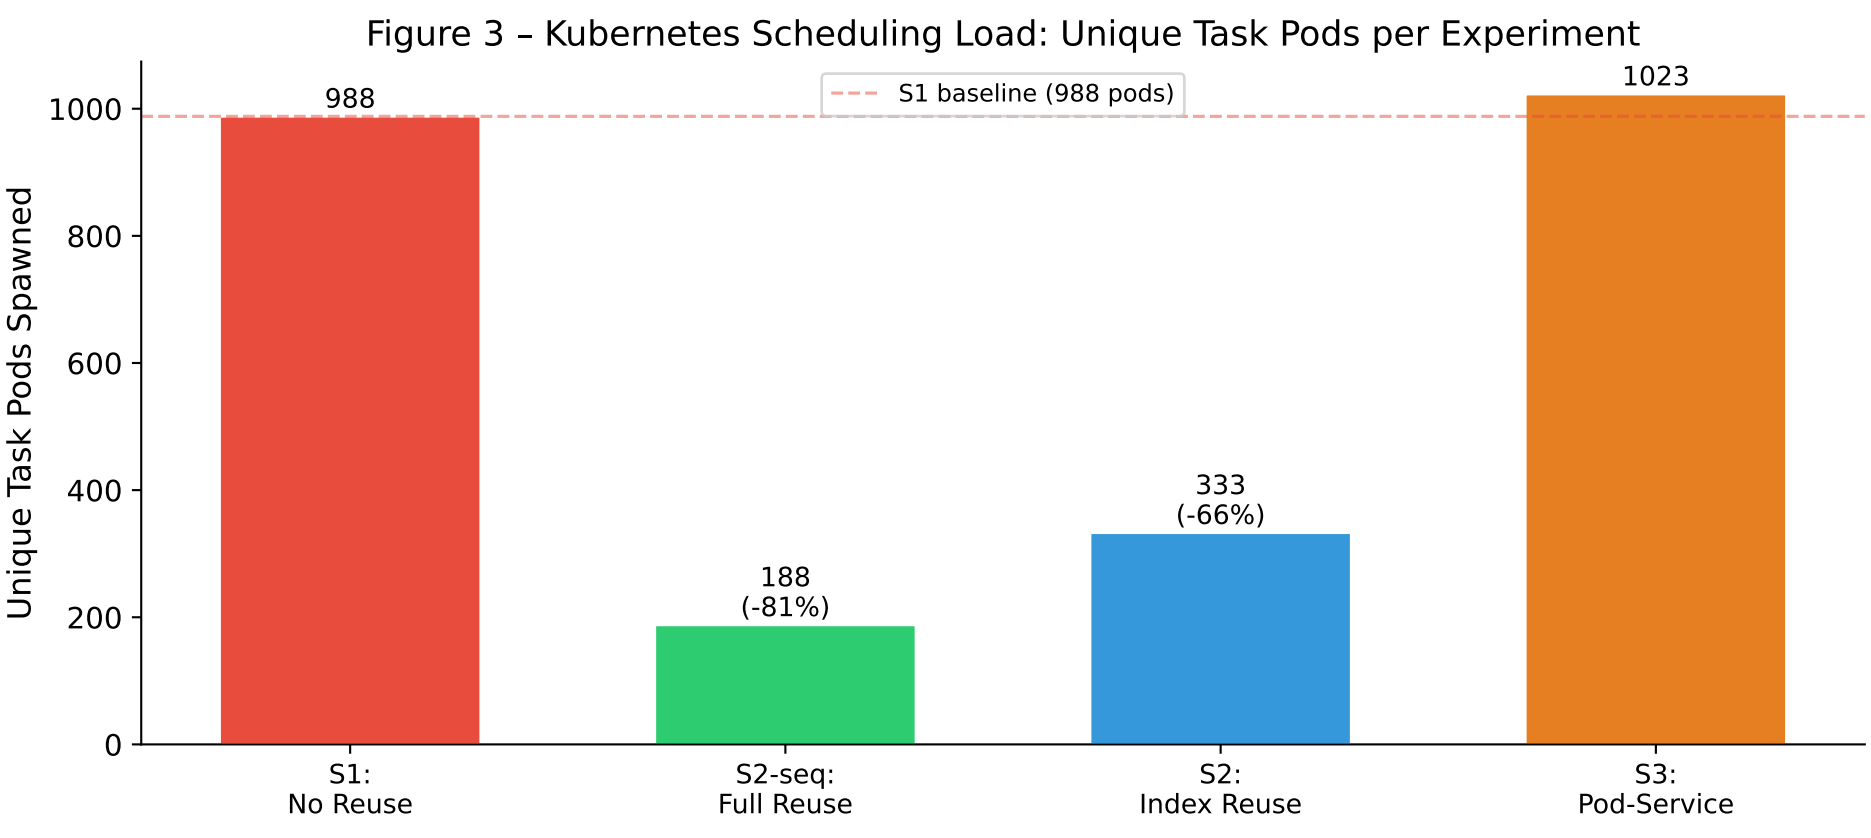
Figure 3 – Kubernetes Scheduling Load: Unique Task Pods per Experiment
Unique Task Pods Spawned
0
200
400
600
800
1000
S1 baseline (988 pods)
988
188
(-81%)
333
(-66%)
1023
S1:
No Reuse
S2-seq:
Full Reuse
S2:
Index Reuse
S3:
Pod-Service