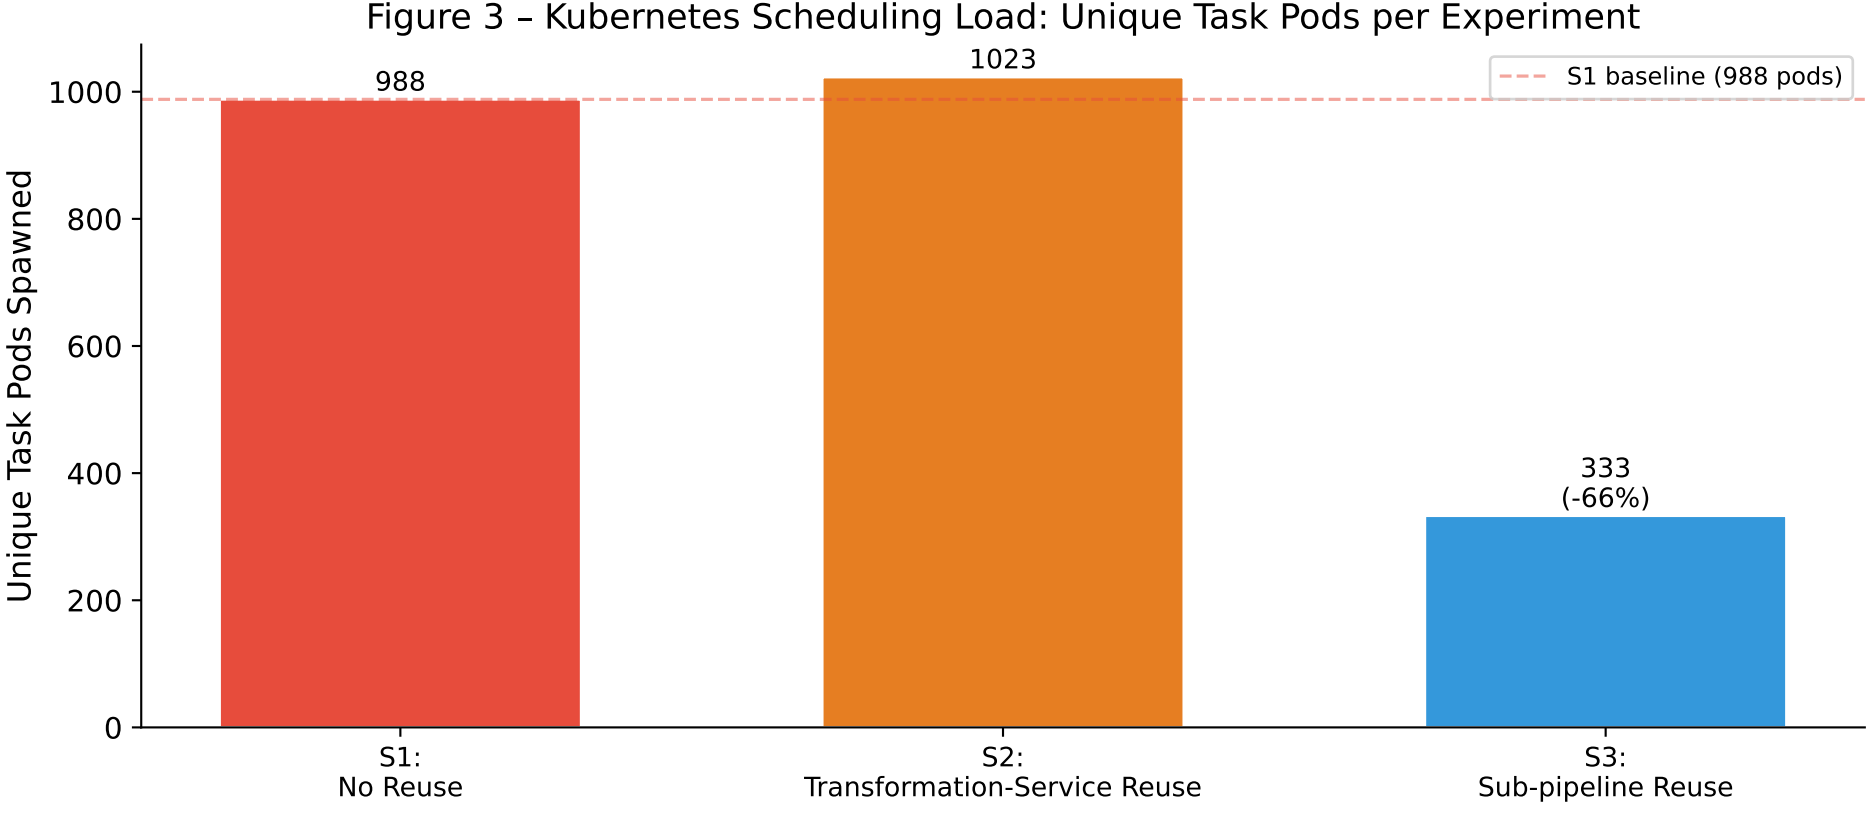
Figure 3 – Kubernetes Scheduling Load: Unique Task Pods per Experiment
Unique Task Pods Spawned
0
200
400
600
800
1000
988
1023
333
(-66%)
S1 baseline (988 pods)
S1:
No Reuse
S2:
Transformation-Service Reuse
S3:
Sub-pipeline Reuse